Regression Line in R

Stat 131A, Fall 2018, Prof. Sanchez

Learning Objectives

- How to run a regression analysis in R
- How to "manually" compute the slope and intercept of regression line
- How to add the regression line to a scatter diagram

Introduction

To make sure we are all on the same page, I should start with a brief discussion about the term "Regression". This term was coined by Sir Francis Galton and introduced in his 1886 paper *Regression Towards Mediocrity in Hereditary Stature.* He used the word "regression" to describe a phenomenon that he observed when analyzing height data of parents and their adult children. What he noticed was that exceptionally tall parents had children who were, on average, less tall; and exceptionally short parents had children who were, on average, less short. Because of the connotations of the word "mediocrity", statisticians later modified the regression statement as: *regression towards the mean.*

For better or for worse, the term regression has evolved and become broader over the years. Nowadays people use the word regression in a more loosely way. The most common terms that you will probably find are "regression analysis" and "regression model" or "regression modeling".

The core idea behind virtually all regression tools is to predict a (quantitative) response variable in terms of one or more predictor variables. In its simplest version, the regression method is used for explaining or modeling the relationship between a single variable Y, called the *response*, *outcome*, *output* or *dependent* variable; and one *predictor*, *input*, *independent* or *explanatory* variable X. This version is commonly referred to as **simple linear regression**, and it is actually the method presented in the FPP book (chapters 10-12)—although it is not explicitly called like that. The reasons why statisticians call it *simple linear regression* are:

- "simple" because there is only one Y and one X
- "linear" because the mathematical model is expressed with a line equation
- "regression" because X is used to predict Y

Note: Pretty much the word "regression" has become synonym for prediction. The word "prediction" implying that the response variable is quantitative. If the response variable to be predicted is of categorical nature, then we talk about "classification". The type of regression we will discuss is the one in which both X and Y are quantitative variables.

Data

To illustrate the regression ideas, we are going to use a data set collected by English mathematician and biostatistician Karl Pearson (1857-1936). Among other things, Pearson was Galton's protégé, he founded the world's first university statistics department at University College London in 1911, and he is considered one of the founding fathers of modern-day Statistics.


Figure 1: Karl Pearson (source: wikipedia)

The data set is in the csv file `pearson.csv` (in the github repository), and it contains the heights of 1078 fathers, and their adult sons.

```
# assembling the URL of the CSV file
# (otherwise it won't fit within the margins of this document)
repo = 'https://raw.githubusercontent.com/ucb-introstat/introstat-spring-2017/'
datafile = 'master/data/pearson.csv'
url = paste0(repo, datafile)

# read in data set
dat = read.csv(url)
```

If you are having problems importing the data from github, you can try this other source:

```
# another way to read in the data
dat = read.csv('http://www.math.uah.edu/stat/data/Pearson.csv')
```

Univariate Exploration

The data frame `dat` contains 1078 rows, and 2 columns.

- `Father`: height of the father (in inches)
- `Son`: height of the son (in inches)

As it is customary, the first thing when analyzing data (one variable at a time), is to obtain summary statistics and look at the distributions:

```r
# basic summary statistics
summary(dat)
```

```
##       Father           Son
##  Min.   :59.00   Min.   :58.50
##  1st Qu.:65.80   1st Qu.:66.90
##  Median :67.80   Median :68.60
##  Mean   :67.69   Mean   :68.68
##  3rd Qu.:69.60   3rd Qu.:70.50
##  Max.   :75.40   Max.   :78.40
```

```r
# number of rows
n = nrow(dat)

# SD of Father
sqrt((n-1)/n) * sd(dat$Father)
```

```
## [1] 2.744553
```

```r
# SD of Son
sqrt((n-1)/n) * sd(dat$Son)
```

```
## [1] 2.814888
```

Histograms

For convenience purposes, to plot histograms for the `Father` and `Son` heights, we can define a vector of class intervals or bins for the argument `breaks` inside the `hist()` function. Looking at the output of `summary()`, we can take a minimum of 55, and a maximum of 80 to create a vector of `bins` with a numeric sequence like this:

```r
bins = seq(from = 55, to = 80, by = 1)
```

Having defined `bins`, we can then graph the histograms:

```r
hist(dat$Father, breaks = bins, las = 1, col = 'gray80')
```

Histogram of dat$Father


```
hist(dat$Son, breaks = bins, las = 1, col = 'gray80')
```

Histogram of dat$Son

Both histograms have nice symmetric bell-shaped distributions. Keep in mind that these are classic textbook examples of variables that follow the normal curve.

Scatter Plot

The next step is to get a picture of the relationship between the `Father` and `Son` heights. The quickest option is to create a scatter diagram with the function `plot()`:

```
plot(dat$Father, dat$Son)
```

To obtain a better visual display, you can modify the point character `pch` value, and add a color with some transparency using hexadecimal notation:

```r
plot(dat$Father, dat$Son, pch = 19, col = '#7396D655')
```


If you want to get a scatter diagram like the one displayed in the textbook (see page 120), then run the following lines of code:

```r
plot.new()
plot.window(xlim = c(58, 80), ylim = c(58, 80))
```

```
points(dat$Father, dat$Son, pch = 20, cex = 0.5)
axis(side = 1, pos = 58, at = seq(58, 80, 2))
axis(side = 2, las = 1, pos = 58, at = seq(58, 80, 2))
```


The way in which this plot is constructed is a bit special. This approach to create graphs in R uses only low-level functions that give you total control of the appearance of graphical elements:

- `plot.new()` starts a new plot frame
- `plot.window()` is used to set up the coordinates of the axes
- `points()` is used to actually plot the dots
- `axis()` is used to plot both the x-axis and the y-axis.

The most important thing to pay attention to when inspecting the scatter diagram is to check whether the cloud of points follows a linear pattern. When this is the case, it makes sense to use of a line to summarize the relationship between the analyzed variables.

Regression Method

The goal is to find a line that fits the data well. In other words, we want to find a line that is as close as possible to all the data points. The typical algebraic equation of a line is $y = mx + b$, where m represents the slope, and b represents the intercept. However, when working within a regression framework, we use a slightly different notation for the equation of a line. The most standard notation is:

$$Y = \beta_0 + \beta_1 X + \epsilon$$

6

where β_0 and β_1 are two unknown constants that represent the intercept and the slope of a line, respectively. In turn, ϵ is the error term that accounts for the spread in the cloud of points that is not captured by the model. Together, β_0 and β_1 are known as the model *coefficients* or *parameters*.

Often, you will find an alternative represention of the model as:

$$Y \approx \beta_0 + \beta_1 X$$

you migth read "$\approx$" as *"is approximately modeled as."* This is a very important detail: we are assuming that there is approximately a linear relationship between X and Y.

We will sometimes describe the equation $Y \approx \beta_0 + \beta_1 X$ by saying that we are *regressing Y on X* (or *Y onto X*).

In our working example, X represents the height of parents, `Father`, and Y represents height of sons, `Son`. Then we can regress `Son` onto `Father` by fitting the model:

$$\texttt{Son} \approx \beta_0 + \beta_1 \texttt{Father}$$

Once we have used the data to produce estimates b_0 and b_1 for the model coefficients, we can predict the heights of sons based on a particular value of a father's height by computing:

$$\hat{y} = b_0 + b_1 x$$

Here we use a *hat* symbol (i.e. the caret), ^, to denote the predicted value. That is, $\hat{y}$ indicates a prediction of Y on the basis of $X = x$.

Estimating the Coefficients

In practice, β_0 and β_1 are unknown. To find estimates b_0 and b_1, we must use the data in X and Y. In other words, we want to find an intercept b_0 and a slope b_1 such that the resulting line is as close as possible to all the data points. There are a number of ways of measuring *closeness*. However, by far the most common approach involves minimizing the *least squares* criterion.

Using the formulas of the regression method, the slope for the regression line is given by:

$$slope = r \times \frac{SD_y}{SD_x}$$

In turn, the formula of the intercept is:

$$intercept = avg_y - slope \times avg_x$$

Using the more standard notation b_0 (intercept) and b_1 (slope), we have that:

$$b_1 = r\frac{SD_y}{SD_x}$$

and

$$b_0 = avg_y - b_1 \times avg_x$$

First let's get the main ingredients:

```r
# number of values (to be used for correcting SD+)
n = nrow(dat)

# averages
avg_x = mean(dat$Father)
avg_y = mean(dat$Son)

# SD (corrected SD+)
sd_x = sqrt((n-1)/n) * sd(dat$Father)
sd_y = sqrt((n-1)/n) * sd(dat$Son)

# correlation coefficient
r = cor(dat$Father, dat$Son)
```

Now let's compute the slope and intercept, and compare them with `reg$coefficients`

```r
# slope
b1 = r * (sd_y / sd_x)
b1
```

```
## [1] 0.5140059
```

```r
# intercept
b0 = avg_y - (b1 * avg_x)
b0
```

```
## [1] 33.8928
```

## Plotting the Regression Line

Having obtained the slope and the intercept estimates, we can use them to get a scatterplot with the regression line on it. The simplest way to achieve this visualization is to first create a scatter diagram with `plot()`, and then add the regression line with the function `abline()`; here's the code in R:

```
# scatterplot with regression line (option 1)
plot(dat$Father, dat$Son)
abline(a = b0, b = b1, lwd = 2)
```


Similarly, you can use the argument `coef` to pass a vector with the intercept and slope values:

```
# scatterplot with regression line (option 2)
plot(dat$Father, dat$Son)
abline(coef = c(b0, b1), lwd = 2)
```


The function `abline()` allows you to add lines to a `plot()`. The good news is that it recognizes objects of class `"lm"`, and when invoked after a `plot()` is created, it will add the regression line. Again, to get a nicer plot, you can do something like this:

```
# scatterplot with regression line
plot.new()
plot.window(xlim = c(58, 80), ylim = c(58, 80))
points(dat$Father, dat$Son, pch = 20, cex = 0.5)
abline(coef = c(b0, b1), col = "tomato", lwd = 2)   # regression line
axis(side = 1, pos = 58, at = seq(58, 80, 2))
axis(side = 2, las = 1, pos = 58, at = seq(58, 80, 2))
```